

Classifying Melanoma Types from Skin Lesion Images

Submitted by:

Nivetha Vedanarayanan


Radhika Rajeevan

Weijun Zhu

Xinran Li


Neural Networks & Deep Learning Project

Advised By : Dr. Farid Alizadeh



Objective

- In this project, we aim to explore the application of Neural Networks in Image Classification for medical diagnosis.
- We are using dermoscopic images of pigmented skin lesions and identifying the melanoma type that it belongs to.



Motivation

- Melanoma are cancerous in nature; however the degree of malignancy differs based on each type.
- In 2015 there were a total of 350,000 cases of skin cancer with 60,000 deaths (17%)
- If diagnosed early, melanoma survival exceeds 95%

Ref- <https://challenge2018.isic-archive.com/>

Introduction



Data source



**Exploratory Data
Analysis**



Pre-processing




Approaches

Building CNN
Feature Extraction
Pretrained Model Architecture
Reusing Pretrained Weights




Comparing Models



Data Source & Extraction

- The Skin Lesion Images were made available on Kaggle as part of a Melanoma Detection Challenge by the International Skin Imaging Collaboration (ISIC).
- There are 10,015 dermoscopic colored images classified into 7 different melanoma types.
- We have extracted the data from Kaggle into Google drive using Kaggle API.

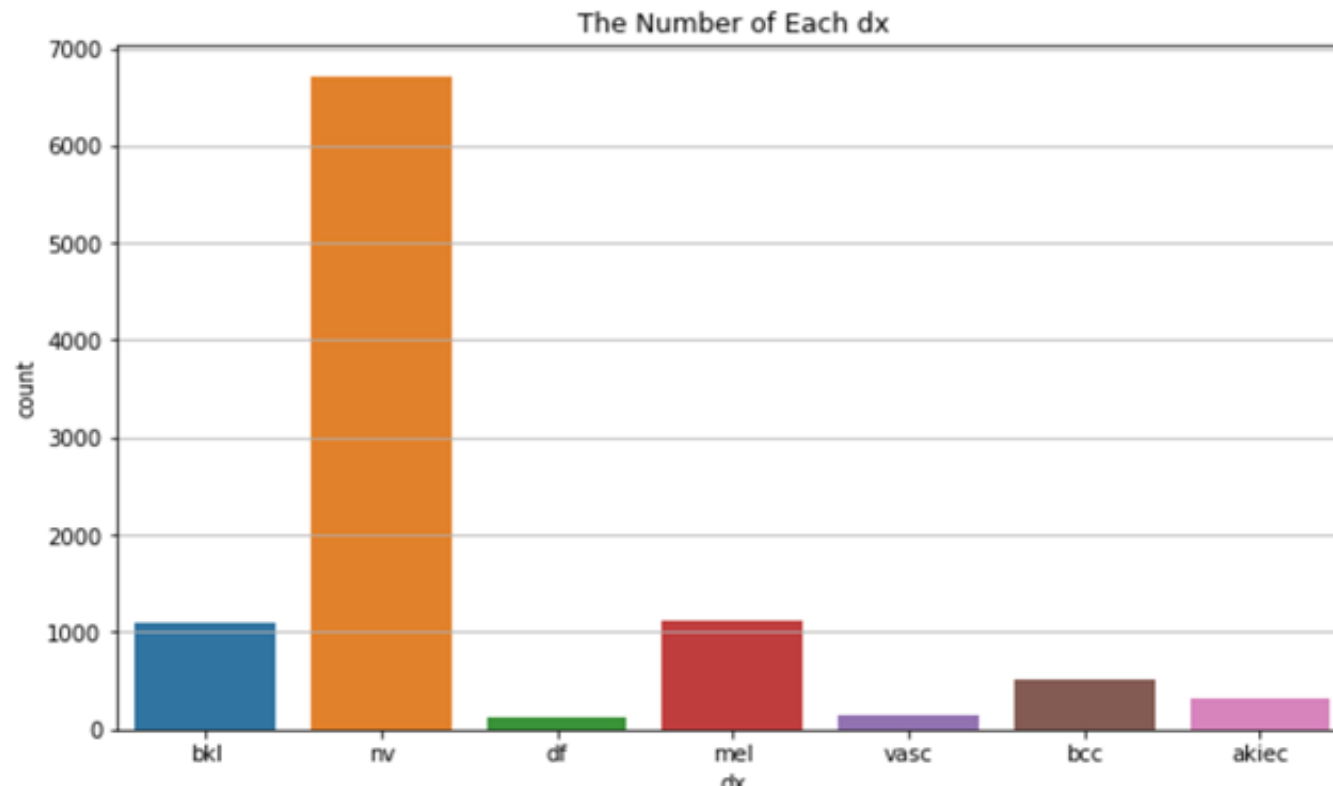


EDA Findings

- The Skin Lesion Images were made available on Kaggle as part of a Melanoma Detection Challenge by the International Skin Imaging Collaboration (ISIC).
- There are 10,015 dermoscopic colored images classified into 7 different melanoma types.
- We have extracted the data from Kaggle into Google drive using Kaggle API.

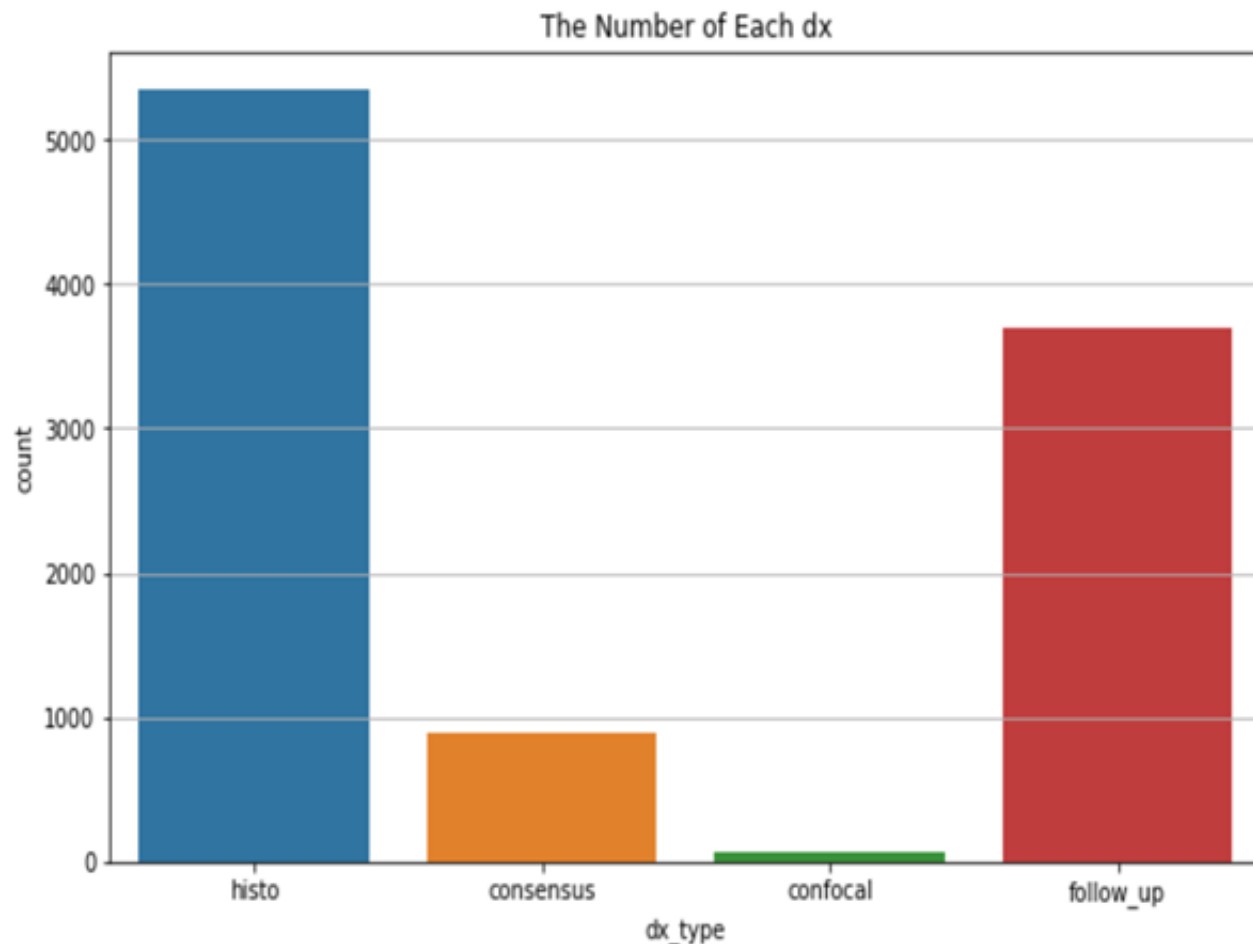
df feature

```
nv      6705
mel     1113
bkl     1099
bcc      514
akiec    327
vasc     142
df       115
Name: dx, dtype: int64
```



- By visualizing data, we can tell that the “nv” type is the most common melanoma type. The quantity of “nv” exceed over $\frac{1}{2}$ of the overall collected images.
- The other melanoma types are: “bkl”, “df”, “mel”, “vasc”, “bcc”, “akiec”

```
histo      5340
follow_up  3704
consensus   902
confocal    69
Name: dx_type, dtype: int64
```




dx type feature


- We notice from the visualization that over 50% of lesions are confirmed through “histo”.
- While “follow_up” play another major role in confirming cases, we could also confirmed by “consensus” and “confocal”.

Some other features... ..


The Density Plot of the Age Feature



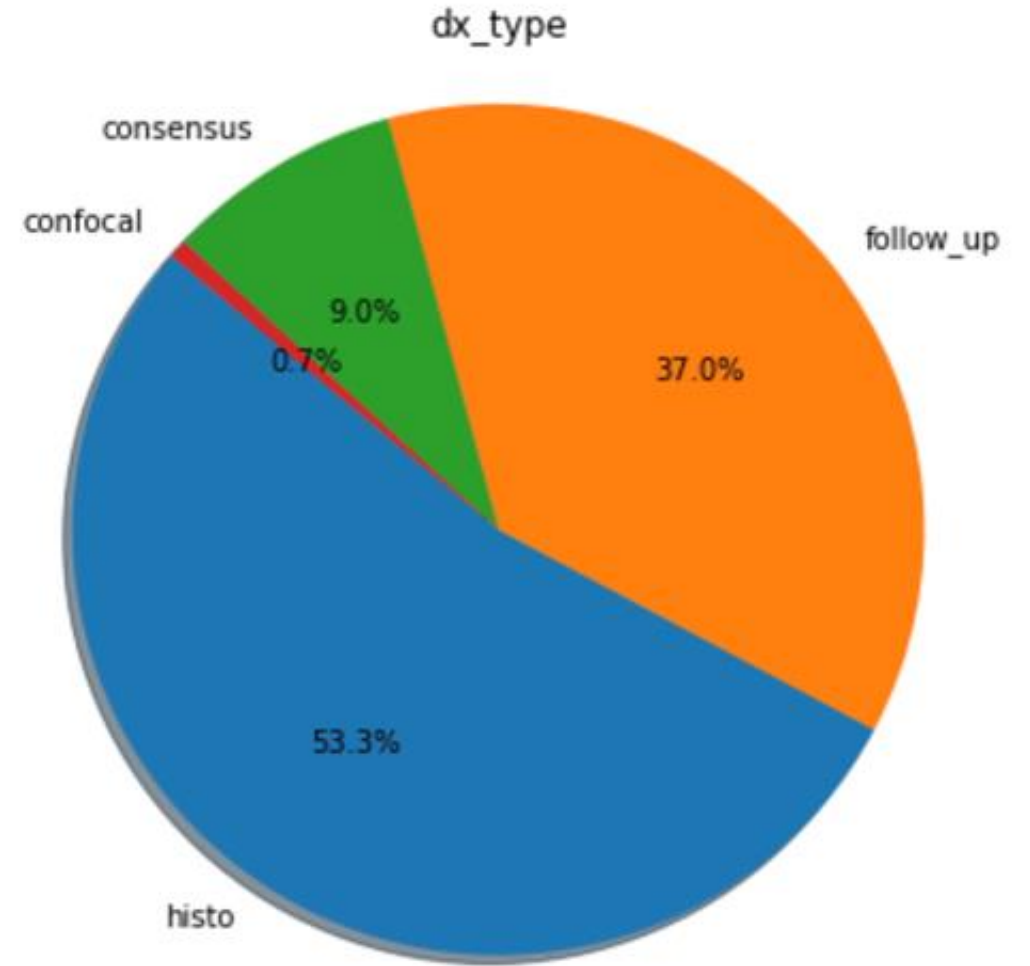

The Number of Each Sex




The Number of Each Location



Another view for the percentage between dx and dx_type






Relationship
between
dx and dx_type
(How & Where)


```
sizes = [5340,3704,902,69]
sizes_dx = [6705,1113,1099,514,327,142,115]
labels = metadata["dx_type"].unique()

plt.figure(figsize=(9,6))
plt.pie(sizes, labels=labels, startangle=90,frame=True)
plt.pie(sizes_dx,radius=0.75,startangle=90)
centre_circle = plt.Circle((0,0),0.5,color='black', fc='white',linewidth=0)
fig = plt.gcf()
fig.gca().add_artist(centre_circle)
```




Preprocessing

- Scaling Data
 - Images were converted to numeric arrays and scaled between 0 and 1 and saved as npy files.
 - We have tried models with 3 different dimensions:
 - $128 \times 128 \times 3$
 - $150 \times 150 \times 3$
 - $224 \times 224 \times 3$
- Class Imbalance
 - Since class proportions are highly skewed, we have used Data Augmentation to generate rotated images for classes with less training data for a class-wise balance.



MobileNet

- MobileNet with 128 x 128 x 3 arrays
 - Train-Val-Test Split
 - Model Architecture
 - Learning Rate
 - Accuracy & Loss
 - Performance on Test Data
 - Confusion Matrix
 - Precision & Recall



MobileNet Architecture


- Freeze the pretrained MobileNet weights

Model: "mobilenet_1.00_128"

Layer (type)	Output Shape	Param #
input_5 (InputLayer)	(None, 128, 128, 3)	0
conv1_pad (ZeroPadding2D)	(None, 129, 129, 3)	0
conv1 (Conv2D)	(None, 64, 64, 32)	864
conv1_bn (BatchNormalization)	(None, 64, 64, 32)	128
conv1_relu (ReLU)	(None, 64, 64, 32)	0
conv_dw_1 (DepthwiseConv2D)	(None, 64, 64, 32)	288
conv_dw_1_bn (BatchNormaliza)	(None, 64, 64, 32)	128
conv_dw_1_relu (ReLU)	(None, 64, 64, 32)	0
conv_pw_1 (Conv2D)	(None, 64, 64, 64)	2048
conv_pw_1_bn (BatchNormaliza)	(None, 64, 64, 64)	256
conv_pw_1_relu (ReLU)	(None, 64, 64, 64)	0

```
print('Total weights : ',len(model_mobilenet.weights))
print('Trainable weights : ',len(model_mobilenet.trainable_weights))
```


```
Total weights : 143
Trainable weights : 8
```



MobileNet Performance

RBS - MITA - Spring 2020


- Learning Rate of 0.0001
- Optimizer - RMSProp



- The validation loss is > 1 and does not drop
- Validation accuracy does not improve beyond 80%

MobileNet Test Data

<matplotlib.image.AxesImage at 0x7




```
# test accuracy
test_loss, test_acc = model_mobilenet.evaluate(mat_data_test,mat_label_test)
print('Test Loss : ',test_loss,' Test Acc : ',test_acc)
```

```
1002/1002 [=====] - 3s 3ms/step
Test Loss : 0.9709897623239401 Test Acc : 0.7614770531654358
```


	Class	Acc		akiec	bcc	bkl	df	nv	vasc	mel
0	akiec	0.424242	akiec	14	2	10	0	5	0	2
1	bcc	0.450980	bcc	6	23	5	0	16	0	1
2	bkl	0.518182	bkl	8	1	57	0	39	1	4
3	df	0.083333	df	0	0	3	1	8	0	0
4	nv	0.953800	nv	3	3	17	0	640	3	5
5	vasc	0.428571	vasc	0	0	2	0	6	6	0
6	mel	0.198198	mel	6	3	14	1	64	1	22

- Identification of 'df' is extremely low : 8%
- Except 'nv' rest of classes have an accuracy of $\leq 50\%$



MobileNet with Data Aug

- MobileNet with Data Augmentation
 - Images are Generated on Imbalanced class images
 - Model Architecture
 - Learning Rate
 - Accuracy & Loss
 - Performance on Test Data
 - Confusion Matrix
 - Precision & Recall



MobileNet-DA Architecture


- Freeze the pretrained MobileNet weights

Model: "sequential_1"

Layer (type)	Output Shape	Param #
mobilenet_1.00_128 (Model)	(None, 4, 4, 1024)	3228864
flatten_1 (Flatten)	(None, 16384)	0
dropout_1 (Dropout)	(None, 16384)	0
dense_1 (Dense)	(None, 1024)	16778240
dropout_2 (Dropout)	(None, 1024)	0
dense_2 (Dense)	(None, 128)	131200
dropout_3 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 64)	8256
dense_4 (Dense)	(None, 7)	455
Total params: 20,147,015		
Trainable params: 20,125,127		
Non-trainable params: 21,888		


```
[ ] print('Total weights : ',len(model_mobilenet.weights))
    print('Trainable weights : ',len(model_mobilenet.trainable_weights))
    print('Non Trainable weights : ',len(model_mobilenet.non_trainable_weights))
```

Total weights : 143
Trainable weights : 89
Non Trainable weights : 54



MobileNet-DA Performance


- Learning Rate of 0.0001
- Optimizer – RMSProp
- $\text{steps_per_epoch} = (\text{train_size} / \text{batch_size})$



- High fluctuation in validation loss
- Validation accuracy improves slightly beyond 80%

MobileNet-DA Test Data

<matplotlib.image.AxesImage at 0>



RBS - MITA - Spring 2020

```
# test accuracy
test_loss, test_acc = model_mobilenet.evaluate_generator(test_generator, steps=50)
print('Test Loss : ', test_loss, ' Test Acc : ', test_acc)
```

Test Loss : 2.1438980102539062 Test Acc : 0.8329097628593445


	Class	Acc		akiec	bcc	bkl	df	nv	vasc	mel
0	akiec	0.133333	akiec	4	4	7	3	11	1	0
1	bcc	0.586957	bcc	3	27	7	2	4	3	0
2	bkl	0.636364	bkl	0	0	63	4	7	25	0
3	df	0.800000	df	0	0	2	8	0	0	0
4	nv	0.630000	nv	0	1	9	0	63	27	0
5	vasc	0.953642	vasc	0	1	12	1	14	576	0
6	mel	0.538462	mel	0	0	0	0	3	3	7

- Identification of 'df' is extremely low : 8%
- Except 'nv' rest of classes have an accuracy of $\leq 50\%$

Downgrades

- Accuracy for 'akiec' has dropped from 42% to 13%
- 'nv' has dropped from 95% to 63%

<matplotlib.image.AxesImage at 0x7




	akiec	bcc	bkl	df	nv	vasc	mel
akiec	14	2	10	0	5	0	2
bcc	6	23	5	0	16	0	1
bkl	8	1	57	0	39	1	4
df	0	0	3	1	8	0	0
nv	3	3	17	0	640	3	5
vasc	0	0	2	0	6	6	0
mel	6	3	14	1	64	1	22

	Class	Acc
0	akiec	0.424242
1	bcc	0.450980
2	bkl	0.518182
3	df	0.083333
4	nv	0.953800
5	vasc	0.428571
6	mel	0.198198

Improvements


- Identification of 'df' has improved from 8% to 80%
- Fewer misclassifications

<matplotlib.image.AxesImage at 0>



	Class	Acc
0	akiec	0.133333
1	bcc	0.586957
2	bkl	0.636364
3	df	0.800000
4	nv	0.630000
5	vasc	0.953642
6	mel	0.538462

	akiec	bcc	bkl	df	nv	vasc	mel
akiec	4	4	7	3	11	1	0
bcc	3	27	7	2	4	3	0
bkl	0	0	63	4	7	25	0
df	0	0	2	8	0	0	0
nv	0	1	9	0	63	27	0
vasc	0	1	12	1	14	576	0
mel	0	0	0	0	3	3	7



MobileNet-DA with Balanced Class

- MobileNet with Data Augmentation
 - Images are Generated on Imbalanced class images
 - Model Architecture
 - Learning Rate
 - Accuracy & Loss
 - Performance on Test Data
 - Confusion Matrix
 - Precision & Recall



PART IV

ResNet50

Weijun Zhu

Residual Neural Network



Deeper neural networks are more difficult to train, and sometimes researchers find the deeper neural network get the worse results than the shallower one.



ResNet was found by Kaimin in 2015, and it is good to fit the model in the deeper neural network.




They explicitly reformulate the layers as learning residual functions with reference to the layer inputs, instead of learning unreferenced functions.

Residual Neural Network

```
import torch.nn as nn
import torch
from torch.nn.init import kaiming_normal, constant

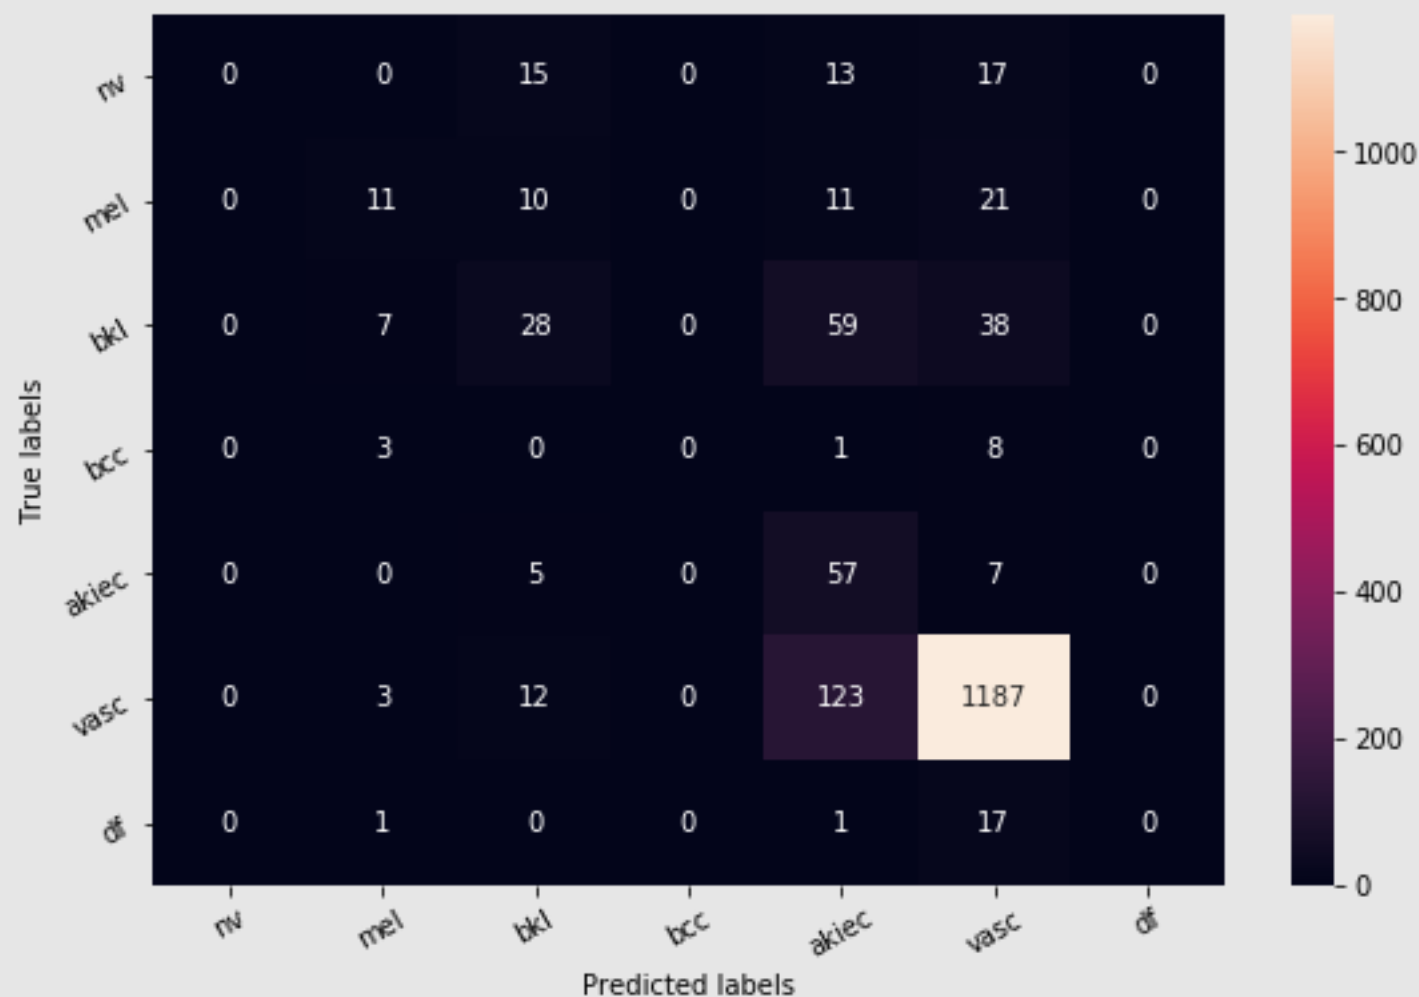
class BasicConvResBlock(nn.Module):
    def __init__(self, input_dim=128, n_filters=256, kernel_size=3, padding=1,
                 stride=1, shortcut=False, downsample=None):
        super(BasicConvResBlock, self).__init__()
        self.downsample = downsample
        self.shortcut = shortcut
        self.conv1 = nn.Conv1d(input_dim, n_filters, kernel_size=kernel_size,
                                padding=padding, stride=stride)
        self.bn1 = nn.BatchNorm1d(n_filters)
        self.relu = nn.ReLU()
        self.conv2 = nn.Conv1d(n_filters, n_filters, kernel_size=kernel_size,
                                padding=padding, stride=stride)
        self.bn2 = nn.BatchNorm1d(n_filters)
    def forward(self, x):
        residual = x
        out = self.conv1(x)
        out = self.bn1(out)
        out = self.relu(out)

        out = self.conv2(out)
        out = self.bn2(out)
        out += residual
        out = self.relu(out)
        return out
```



ResNet 50


Confusion Matrix



	precision	recall	f1-score	support
akiec	0.00	0.00	0.00	45
bcc	0.44	0.21	0.28	53
bkl	0.40	0.21	0.28	132
df	0.00	0.00	0.00	12
mel	0.22	0.83	0.34	69
nv	0.92	0.90	0.91	1325
vasc	0.00	0.00	0.00	19
accuracy			0.78	1655
macro avg	0.28	0.31	0.26	1655
weighted avg	0.79	0.78	0.77	1655



MobileNet

Confusion Matrix




	precision	recall	f1-score	support
akiec	0.55	0.13	0.21	45
bcc	0.80	0.15	0.25	53
bkl	0.82	0.07	0.13	132
df	0.08	0.08	0.08	12
mel	0.16	0.48	0.24	69
nv	0.90	0.95	0.93	1325
vasc	0.88	0.37	0.52	19
accuracy			0.80	1655
macro avg	0.60	0.32	0.34	1655
weighted avg	0.84	0.80	0.78	1655

ResNet 50; epoch=10




MobileNet; epoch=30

Training and validation loss



Training and validation cat accuracy



Reference

1. Deep Residual Learning for Image Recognition, Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun
2. Identity Mappings in Deep Residual Networks, Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun
3. MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications, Andrew G. Joward, Menglong Zhu, Bo Chen, Dmitry Kalenichenko, Weijun Wang, Tobias Weyand, Marco Andreetto, Hartwig Adam
4. <https://www.zhihu.com/search?type=content&q=%E6%AE%8B%E5%B7%AE%E7%A5%9E%E7%BB%8F%E7%BD%91%E7%BB%9C>
5. <https://towardsdatascience.com/introduction-to-resnets-c0a830a288a4>



THANK YOU